Analysis of city council meeting minutes

Institut des Sciences du Digital, Management et Cognition
Université de Lorraine

November 29, 2024

Oyetunji Abioye, Alberto Lorente Galé, Mina Oulhen and Ziyan Xu

Project Overview

Datapolitics Project

- **Objective**: Develop an automated detector to identify and categorize projects implemented by local authorities.
- **Scope**: Examine around **20,000** geothermal energy PDF documents from the past five years. Ensure the methodology can be applied to various local projects

Data Overview

- **doc_id**: Unique identifier for each document.
- **url**: Original source URL of the document.
- **cache**: Link to the cached PDF version.
- **fulltext**: Link to the plain text version of the document.
- **nature**: Automatically classified document type (e.g., deliberation, minutes).
- **published**: Publication date of the document.
- **entity_name**: Name of the local authority responsible for the document.
- **entity_type**: Type of the entity (e.g., municipality, intercommunality).
- **geo_path**: Hierarchical administrative path indicating the geographical scope.

Tasks

Filtering and Classification Process

1. **First Level: Binary Filter**
   - **Concerns a Geothermal Project**
   - **Unrelated to a Geothermal Project**
2. **Second Level: Project Stages**
   - Idea/Wish
   - Preliminary Studies
   - Budget Voted for the Definitive Project
   - Implementation in Progress
   - Implementation Completed
3. **Final Level: Data Extraction**
   - Initial Budget
   - Final Cost
   - Estimated Duration
   - Actual Duration


WHEN YOU HAVE TO ANNOTATE 20K
DOCS BUT YOU'RE JUST A CHILL GROUP

Exhausted options to deal with the annotation process:

- Generating the annotation via clustering with LLM Embeddings [1].

Dataset	Embed.	Best Alg.	F1S	ARI	HS	SS	CHI	Total
DS1	TF-IDF	k-means	0.67	0.38	0.46	0.016	4	0/5
	BERT	Spectral	**0.85**	**0.60**	0.63	**0.118**	25	3/5
	OpenAI	k-means	0.84	0.59	**0.64**	0.066	13	1/5
	LLaMA-2	k-means	0.41	0.09	0.17	0.112	**49**	1/5
	Falcon	k-means	0.74	0.39	0.48	0.111	34	0/5
DS2	TF-IDF	Spectral	0.82	0.63	0.58	0.028	8	0/5
	BERT	AHC	0.74	0.58	0.53	0.152	37	0/5
	OpenAI	AHC	**0.90**	**0.79**	**0.75**	0.070	19	3/5
	LLaMA-2	k-means	0.51	0.21	0.25	0.137	69	0/5
	Falcon	k-means++	0.45	0.26	0.30	**0.170**	**85**	2/5
DS3	TF-IDF	Spectral	0.35	0.13	0.28	-0.002	37	0/5
	BERT	k-means	0.43	0.25	0.44	0.048	412	0/5
	OpenAI	k-means	**0.69**	**0.52**	**0.66**	0.035	213	3/5
	LLaMA-2	AHC	0.17	0.11	0.26	0.025	264	0/5
	Falcon	k-means	0.26	0.15	0.30	**0.071**	**1120**	2/5
DS4	TF-IDF	k-means	0.29	0.13	0.48	0.034	17	0/5
	BERT	k-means	0.35	0.24	0.55	**0.072**	61	1/5
	OpenAI	k-means	**0.38**	**0.26**	**0.58**	0.053	42	3/5
	LLaMA-2	k-means	0.21	0.11	0.40	0.053	88	0/5
	Falcon	k-means++	0.27	0.16	0.48	0.071	**92**	1/5

Results from Petukhova et al. (2024)

Educated (?) wish: the embedding distance captures the project vs non-project relationship.
Process for cluster explainability: 2 clusters for each document type. Train a RF model with TF-IDF transformations of the text. Look at the most important features.
Top 20 most important words for different document types:

- **acte.arrete**: région, direction, subdélégation, recueil, agrément, relatif, mission, transport, arrêter, mois, département, décret, national, devoir, e, être, signer, maire, faire, délai.
- **bdj**: local, avoir, fonds, travail, investissement, compte, service, être, opération, personnel, subvention, public, recette, général, année, entrer, devoir, dépense, budget, charge.

inconclusive results = prompt LLAMA

RAG

After setting up a RAG *poc* and doing some trials with Llama 3.2 1B the retrieval was not as accurate as we hoped. Example with one document loaded with *sentence-camembert-base* embeddings:

US: C'est un project?

LLAMA: Il y a des projets dans l'extrait du registre des délibérations du Conseil municipal le 1er juillet 2022.

Node

Mis en ligne le : 12/07/2022 - vUle t1eV n1ers.-le.-bel Séance du Conseil Municipal du vendredi I juillet 2022 Extrait du Registre des Délibérations du Conseil Municipal Ville de Villiers-le-bel Séance ordinaire du vendredi 1 juillet 2022 N°1/Com pte rendu Compte rendu commenté (procès-verbal) de la séance du Conseil Municipal du 24 mai 2022 Le vendredi.

Chain-of-Thought Prompting

So we are using Chain-of-Thought (CoT) prompting to guide a GPT4-o model towards structured, logical answers. But its not working that well either:

US: You are a helpful AI Question Answerer. You are about to be passed a project description text in French by a User. Your task is to extract the estimated duration of the project. Make sure to generate a series of logical steps and reason about the problem before generating a solution. (...)

GPT: The actual duration is: 3. Chain-of-Thought: (...) in the context here of understanding a duration from a publishing perspective, the publication covered listings of 3 pages regarding a decision from 25th August. Hence, the most visible duration specifiable from **document length is 3 pages**, interpreting via the listings as grasped directly from the text.

Conclusion

Moving forward:

- Prompt LLMs to annotate the target variable for the binary classifier.
- Keep working on the RAG system with CoT generation.

References

 A. Petukhova, J. P. Matos-Carvalho & N. Fachada – "Text clustering with llm embeddings", 2024.